

Voice Recognition with Topological Data Analysis

Joy Patel, Gautam Hathi, Austin Hua

May 4, 2016

Contents

1	Introduction	3
2	Data	3
3	Process	3
3.1	Audio Feature Extraction for Topology	4
3.2	Distance Metrics	5
3.3	Statistics and Machine Learning	5
3.3.1	Two Sample T Test	6
3.3.2	Machine Learning Classifiers	6
4	Results	6
4.1	Two Sample T Test Results	6
4.2	Classification Results	7
4.3	Discussion	7
5	Conclusion	8
6	Bibliography	9

1 Introduction

Our project focuses on the topic of voice discrimination. While the problem of voice recognition—which involves converting voice to text—is widely discussed and explored, the problem of distinguishing between speakers—which we call voice discrimination—is also important. In particular, there are lots of situations involving new devices such as the Amazon Echo or Google Home where it is important not only to know what is being said but who is speaking. Voice discrimination can be a useful tool in these situations. Current voice technology is very single-user centric, but improving voice discrimination capabilities opens up a wide range of multi-user applications. Another possible application is providing another form of biometric security, potentially allowing users to unlock their phones by voice. These applications all hinge on improving the accuracy of voice discrimination.

In this project, we explore a topological approach to voice discrimination. We explore different techniques for processing voice samples that can be used to create good topological features from the voice data. We then attempt to use a number of machine learning and statistical techniques to accomplish voice discrimination from topological features, and compare the results of using topological features vs. non-topological features.

2 Data

For this project, we collected voice samples from 4 different people. We had each of these people say the phrase “open sesame” 40 times. “Open sesame” was chosen because the length of the phrase provided more data and varied sounds to allow for better discriminatory power than a single word might provide. A sample refers to a single instance of a single person saying this phrase. This gave us a dataset with many different people saying the same phrase, and we could then proceed to try and discriminate between samples in the dataset from different people. We could then proceed to try and discriminate between samples in the dataset by the differences in the way people say the same phrase. Some of the ways these differences manifest themselves are in pitch, intonation and inflection, and rhythm and cadence. We aim to take advantage of the effects of these differences by extracting various features as well as using TDA to compute the degree of similarity between samples. In the machine learning statistical context, each person is a class. A voice sample is a signal file, i.e., it can be thought of as a vector in \mathbb{R}^L where L is the time length of the audio file (sample in 44 kHz).

3 Process

We explored three components which we put together into a voice discrimination pipeline.

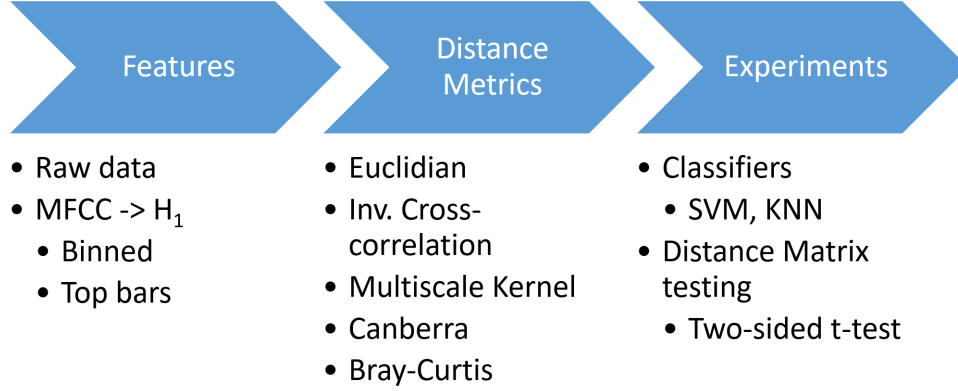


Figure 1: Process Pipeline

3.1 Audio Feature Extraction for Topology

We processed each audio sample using a couple of different feature extraction libraries. To analyze the audio sample topologically, we computed 1st dimensional homology from these features as well as from the raw audio data using the Rips complex. From the resulting persistence diagram, we took the top 10 bars as well as vectors created by binning the persistence diagram.

From the audio data, we compute Mel-Frequency Cepstral coefficient (MFCC) representation which is time series based on a “perceptually motivated log frequency and log power short-time Fourier transform that preserves timbral information” [2] [3]. MFCC features appeared to give us the best homology results. The 3rd MFCC coefficient of an audio sample generally created point clouds with visible cycles in 2D PCA. When we did TDA on the point cloud, we saw clearly distinguished 1D persistence points above the diagonal [3]:




Figure 2: 1D Homology Persistence Diagram of MFCC Features from Voice Sample

For subsequent steps in the process, we used the following features which we generated:

1. Persistence Diagram: 1D persistence diagrams generated from sliding windows over MFCC features (for use with multi-scale kernel) [3]
2. Top 10 Persistence Bars: Top 10 (death - birth) values from 1D persistence diagrams generated from sliding windows over MFCC features (vector in \mathbb{R}^{10})
3. Binned Persistence Diagrams: Binned 1D persistence diagrams generated from sliding windows over MFCC features with 36 bins (vector in \mathbb{R}^{36})
4. Raw Data: Raw audio sample data (vector in \mathbb{R}^t where t is the length of the audio file)

3.2 Distance Metrics

We explored several different distance heuristics and metrics for statistical comparisons and classification. Take note of their inputs, that is, they are metrics and heuristics between different types of features, e.g., integer valued vectors, real valued vectors, signal functions, and persistence diagrams.

1. Real Valued Vector Spaces Euclidean Metric
 $d : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ via $d(u, v) = \sqrt{\sum_{i=1}^n (u_i - v_i)^2}$
2. Integer Valued Vector Spaces Canberra Metric
 $d : \mathbb{Z}^n \times \mathbb{Z}^n \rightarrow \mathbb{R}$ via $d(u, v) = \sum_{i=1}^n \frac{|x_i - y_i|}{|x_i| + |y_i|}$ [?]
3. Integer Valued Vectors Spaces Bray-Curtis Metric
 $d : \mathbb{Z}^n \times \mathbb{Z}^n \rightarrow \mathbb{R}$ via $d(u, v) = \frac{\sum_{i=1}^n |x_i - y_i|}{\sum_{i=1}^n |x_i + y_i|}$ [?] [6]
4. Real Valued Signal Inverse Max Cross Correlation
 $d : \mathbb{C}^1 \times \mathbb{C}^1 \rightarrow \mathbb{R}$ via $d(f, g) = \frac{1}{\max_{\tau \in (-\infty, +\infty)} (\int_{-\infty}^{+\infty} f^*(t)g(t+\tau)dt)}$ where f^* is the conjugate of f [1]
5. Persistence Diagram Distance Multiscale Kernel
 $d : PD \times PD \rightarrow \mathbb{R}$ via $d(P, Q) = \frac{1}{8\pi\sigma} \sum_{p \in P, q \in Q} (\exp(-\frac{\|p-q\|^2}{8\sigma}) - \exp(-\frac{\|p-\bar{q}\|^2}{8\sigma}))$ where $\|v\|$ is the Euclidean Metric and if $q = (a, b)$, then $\bar{q} = (b, a)$ (note: σ is a parameter choice; we use $\sigma = 1.0$) [4]

3.3 Statistics and Machine Learning

In our research, we wish to experiment with both the predictive capabilities of the distance functions with a variety of features. To discriminate between voice samples from different people, we put the features extracted from the samples into statistical tests and classifiers with a correspondingly appropriate distance metrics.

3.3.1 Two Sample T Test

We took sets of feature vectors from voice samples of different people, calculated distances between pairs of vectors within and pairs of vectors across sets, and ran a two-sided t-test to determine whether the distances of vector pairs within a set were statistically distinguishable from the distances of vector pairs across sets.

More concretely: first choose a feature and a corresponding distance metric d . Let C_i be collection of that feature for person i . So, for instance, if the feature you chose was MFCC binned persistence diagrams, then C_i would be the set of MFCC binned persistence diagrams for person i . Let C_j, C_k feature from voice samples from 2 different people (i.e., 2 different classes). For C_k and C_j , construct S_1 and S_2 as follows: $S_1 = \{d(x, y) : x, y \in C_k, x \neq y\}$ and $S_2 = \{d(x, y) : x \in C_k, y \in C_j\}$. Using S_1 and S_2 , run a two sample t-test. The results of this test will show whether the feature and distance metric chosen have discriminatory power.

3.3.2 Machine Learning Classifiers

We trained machine learning classifiers on feature vectors extracted from voice samples of people and used the different distance metrics mentioned above for classification using a 60% training, 40% testing split of the data. We experimented with both binary classification (using 2 labels at a time) as well as multi-label (4 labels) classification using the following sets of (feature, metric) pairs: (Binned Persistence Diagrams, Canberra Metric), (Binned Persistence Diagrams, Bray-Curtis Metric), (Binned Persistence Diagrams, Euclidean Metric), (Top 10 Persistence Bars, Euclidean Metric), (Raw Data, Inverse Cross Correlation), (Persistence Diagram, Multiscale Kernel Metric). For the binary classification, we run each of these (feature, metric) of all pairwise combinations of the people, compute the ratio of $\frac{\text{true positive}}{\text{testing size}}$ per combination, and average the ratios to get an average true positive ratios. Since we use our own metrics for distance computation, we need classifiers that are also based in some space with distance (metrics and heuristics). Thus, our classification experiments are run using SVM and KNN classifiers.

4 Results

4.1 Two Sample T Test Results

The results for a 2 sample T test are shown below for 2 of the classes.

(feature, metric)	Test Statistics	P-Value
Binned Persistence Diagrams, Canberra Metric	-21.420	9.848e-91
Binned Persistence Diagrams, Bray-Curtis Metric	-7.481	1.185e-13
Binned Persistence Diagrams, Euclidean Metric	-4.218	2.601e-05
Top 10 Persistence Bars, Euclidean Metric	2.112	0.0348
Raw Data, Inverse Cross Correlation	-0.337	0.736
Persistence Diagram, Multiscale Kernel Metric	-2.405	0.016

Figure 3: Two Sample T-Test Results

4.2 Classification Results

The results for binary classification (2 labels) are shown below. The table shows the average of the true positive ratios across all pairwise combinations of the classes.

(feature, metric)	SVM	KNN
Binned Persistence Diagrams, Canberra Metric	0.140625	0.8125
Binned Persistence Diagrams, Bray-Curtis Metric	0.21875	0.78125
Binned Persistence Diagrams, Euclidean Metric	0.36458333	0.77604167
Top 10 Persistence Bars, Euclidean Metric	0.48958333	0.53645833
Raw Data, Inverse Cross Correlation	0.42708333	0.71875
Persistence Diagram, Multiscale Kernel Metric	0.57291667	0.5

Figure 4: Binary Classification Average Results

The true positive ratio results for multi-label classification (4 labels) are shown below.

(feature, metric)	SVM	KNN
Binned Persistence Diagrams, Canberra Metric	0.0	0.5625
Binned Persistence Diagrams, Bray-Curtis Metric	0.046875	0.546875
Binned Persistence Diagrams, Euclidean Metric	0.078125	0.515625
Top 10 Persistence Bars, Euclidean Metric	0.265625	0.15625
Raw Data, Inverse Cross Correlation	0.296875	0.59375
Persistence Diagram, Multiscale Kernel Metric	0.3125	0.25

Figure 5: 4 Multi-Label Classification Results

4.3 Discussion

In our discussion of the results, we mainly aim to see how topological features compared with the raw cross-correlation (non-topological) feature. We first conducted the two

sample T-Tests. The two sample T-Test clearly shows significance at a 1% alpha level for the (feature, metric) pairs involving the Binned Persistence Diagrams feature while showing a very high p-value for the Raw Data. This gave us reason to believe that topological tools may be useful for machine learning and are worth investigating.

Afterward, we conducted the machine learning classification experiments, e.g., experiments that may demonstrate how our voice recognition might be used in real world applications. For the binary classification, we found that our best voice discrimination pipeline used (Binned Persistence Diagrams, Canberra Metric) and a KNN classifier with a true positive ratio of 0.8125. Using this pipeline, we were able to achieve better classification results than with (Raw Data, Inverse Cross Correlation) which had a 0.71875 ratio. Our tests also showed that groups of voice samples processed with these features were statistically distinguishable. This means that topological features are promising for voice discrimination purposes with regard to binary classification.

For the multi-label classification, the (Raw Data, Inverse Cross Correlation) with 0.59375 ratio ended up outperforming the topological features. However, the best topological pair was (Binned Persistence Diagrams, Canberra Metric) with a close 0.5625 ratio so even though the Raw Data did better, it was not by a large margin.

We note that the KNN classifier performed better than the SVM classifier. This indicates that the data we fed into the classifiers may not have been linearly separable, even under the different distance metrics we used.

5 Conclusion

Overall, our experiments demonstrate that topological features have distinguishing power for voice data.

There are still some issues to be resolved in future work. First, parameter tuning is an important step in the process that must be explored further. We had to experiment with parameters for the sliding windows in order to ensure that the MFCC feature extraction worked properly. We also note that binning our persistence diagrams worked well with the particular number of bins that we used, and it is unclear how much the number of bins could be varied while still achieving good results.

Performance issues are also a significant consideration. Some of our distance metrics required careful implementation in order to get any reasonable performance, and the implementations we used may not be scalable for large datasets. Determining which distance metrics deliver the best classification results with a scalable implementation would be an important step towards making TDA-based voice discrimination viable.

Finally, while our results are promising, we cannot classify voice samples with near-certainty, as would be desired for real-world applications. Classification accuracy would likely be improved with more data. We had relatively few samples and our models did better with more data. We also note that while we had good results with binary classification, our results were not as good with multi-class classification, indicating that our pipeline may be good for applications such as a voice lock, where the type of discrimination needed is inherently binary (the user in one class, and everyone else in the other). Future works should aim to build a multi-level classifier combining aspects for both raw and topological features.

6 Bibliography

References

- [1] Bracewell, R. *Pentagram Notation for Cross Correlation* The Fourier Transform and Its Applications. New York: McGraw-Hill, pp. 46 and 243, 1965.
- [2] Bruce P Bogert, Michael JR Healy, and John W Tukey. *The quefreny alanalysis of time series for echoes: Cepstrum, pseudo-autocovariance, cross-cepstrum and saphe cracking*. In Proceedings of the symposium on time series analysis, volume 15, 209-243, 1963. Web. 4 May. 2017.
- [3] Chirstopher J Tralie and Paul Bendich. *Cover Song Identification with Timbral Shape Sequences* July 21, 2015. Web. 4 May. 2017. <https://arxiv.org/pdf/1507.05143.pdf>
- [4] Reininghaus Jan, Huber Stefan, Bauer Ulrich, Kwitt Roland. *A Stable Multi-Scale Kernel for Topological Machine Learning* Dec 21, 2014. Web. 4 May. 2017. <https://arxiv.org/pdf/1412.6821.pdf>
- [5] Scikit-learn Developers. *Sklearn neighbors DistanceMetric* Sklearn neighbors DistanceMetric-Scikit-learn Documentation. N.p., 2016. Web. 04 May 2017. <http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.DistanceMetric.html>
- [6] Scipy community. *scipy spatial distance braycurtis* scipy spatial distance braycurtis Documentation. N.p., 2016. Web. 04 May 2017 <https://docs.scipy.org/doc/scipy-0.19.0/reference/generated/scipy.spatial.distance.braycurtis.html>